

A tutorial/walkthrough of basic lasso and ridge model selection techniques

This is a stylized version of a lasso and ridge commands in Stata. The best way to follow along is to download the accompanying “lasso+ridge.do” script which includes on the Stata commands from this file. Be sure to change the file paths to match those of your computer.

Introduction

From what little I know about machine learning (ML) is that Stata has been somewhat slow to the game. Version 16, released less than a year ago in the summer of 2019, finally included an in-house command **lasso** that performs many of the functions that were previously found only in community-written packages. I learned using one of these packages called lassopack – which includes the command **lasso2**. This package is available on SSC and we’ll use that package here to demonstrate lasso and ridge algorithms. I don’t have much experience with the in-house command built into version 16 and the few times I have used it, I get results that I don’t expect. When I have time I’ll include a separate script that directly compares the results from **lasso2** and **lasso**. In short, I’m skeptical of **lasso** because the command doesn’t produce results that I know I should be getting, but I’m very very open to the possibility that I’m not doing something right. If that’s the case, please comment and/or show me how to actually use **lasso** properly.

Getting Started

Anyway, The first thing we need to do is download lassopack from the SSC. We could do so by running the install command from the ssc, which would look like this

```
. ssc install pdslasso
```

However, the authors of this particular package (see <https://statalasso.github.io/>) note that they update their GitHub-hosted version more frequently than that on the SSC. We’ll install directly from GitHub instead and use

```
. net install lassopack, from("https://raw.githubusercontent.com/statalasso/lassopack/master/lassopack_v131/") replace
```

We’ll also be using the estout package to store results from the lasso commands. You probably know about estout but, if not, I highly recommend familiarizing yourself with it as it will become more and more useful to you over time.

```
. ssc install estout, replace
```

Whereas traditional econometrics cares about making sense of relationships between people, actions, and things in the world, the machine-learning approach doesn’t really care about any of that. Instead, machine-learning classifies, predicts, and seeks to build the best model independent of theoretical or historical context. Take Google; they just want to know “yes” or “no”. Will you click on the ad or not? They’ve likely got hundreds of variables on “you” and they’re pretty good at predicting the answer to that question, for each of

us. The World Bank or J-PAL, residing comfortably in the classical econometrical tradition, care much more about making sure the model makes sense in context: does climate change incite violence? Interestingly, from an econometrics perspective, this question becomes very difficult to answer in the affirmative (even though we all know the meteorological and biological evidence of climate change is immense) since very few researchers have told a convincing story that isolates climate change-induced weather variability from other economic or political factors which also affect intermediate factors – such as food or water insecurity – that may incite violence. (A fantastic paper by Selby et al that examines climate change and the Syrian Revolution is linked here: <https://www.sciencedirect.com/science/article/pii/S0962629816301822>) Econometrics cares about “why” and machine-learning cares about what the best model looks like; in a way the former is much harder to do well, in my opinion. But it’s also pretty easy to have machine-learning go horribly wrong. The rest of this tutorial will hopefully set us up with some basic conceptual and coding understanding to, at least, not totally mess up machine-learning.

Traditional Approach

Now we need a dataset. Let’s first use one of Stata’s built-in datasets to demonstrate an instance where ML won’t be as useful.

```
. sysuse auto, clear
```

Now let’s say we want to predict the car’s price. An econometric approach might ask what factors predict price (in miles/gallon) by examining the context of the US auto market in the late 1970s. We might surmise that foreign cars, subject to import tariffs, are likely to be more expensive; also, cars that weigh more might indicate less efficiency. We could run lasso or ridge here, but using our rough “theory of change” we can tell a pretty convincing story already: foreign and weight are two strong predictors using OLS.

```
. reg price          foreign weight
```

Source	SS	df	MS	Number of obs	=	74
Model	316859273	2	158429637	F(2, 71)	=	35.35
Residual	318206123	71	4481776.38	Prob > F	=	0.0000
				R-squared	=	0.4989
				Adj R-squared	=	0.4848
Total	635065396	73	8699525.97	Root MSE	=	2117

price	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
foreign	3637.001	668.583	5.44	0.000	2303.885 4970.118
weight	3.320737	.3958784	8.39	0.000	2.531378 4.110096
_cons	-4942.844	1345.591	-3.67	0.000	-7625.876 -2259.812

Lasso

Great, but what happens if we have a dataset about which we know very little? Let’s take a look at this one about wine quality from the UCI machine learning repository. When you find a dataset you should always keep in mind what the owners allow you to do with it – here we’re ok to use it for research purposes as long as we cite it, so we will: P. Cortez, A. Cerdeira, F. Almeida, T. Matos and J. Reis. Modeling wine preferences by data mining from physicochemical properties. In Decision Support Systems, Elsevier, 47(4):547-553,

2009.

```
. import delimited using "http://archive.ics.uci.edu/ml/machine-learning-databases
> es/wine-quality/winequality-red.csv" , clear
```

Great – now let’s describe dataset:

```
. describe, short
```

Contains data

```
obs:      1,599
```

```
vars:      12
```

Sorted by:

```
Note: Dataset has changed since last saved.
```

Looks like each observation is a wine sample, and we have 12 variables, including **quality** that tells us how “good” the wine is. Let’s see if we can predict quality from the other variables at our disposal. I know nothing about the chemistry of wine, so I can’t use my economic background to help me produce any sort of story. We’ll use lasso to help us with **model selection**: it will tell us which variables to include minimize the sum of the error term.

We’ll use the command called **lasso2** that we downloaded above. We’ll focus on model selection here, but I recommend you check out the package website for all that it can do:

<https://statalasso.github.io/docs/lassopack/>. To see the documentation, type

```
. help lasso2
```

If you look at the help file, you’ll see that we need to tell Stata in order to make the command run: 1) a dependent variable, 2) and a list of candidate explanatory variables from which the algorithm will select the “best fit” model. If we browse, we see that **quality** is the main outcome variable. Let’s create a global called **winevars** to store all of the potential explanatory variables.

```
. global winevars      fixedacidity volatileacidity ///
>                      citricacid residualsugar

.                      chlorides freesulfurdioxide ///
>                      totalsulfurdioxide ph sulphates alcoho
> 1
command chlorides is unrecognized
```

Now let’s run our lasso command. Remember that lasso is an algorithm that selects explanatory variables as a function of lambda, a coefficient “penalty”. The one option worth noting with this command is: `alpha(1)`. This tells Stata to run a **lasso** algorithm as opposed to ridge (see more explanation in the theory file). I’ll include other options that you can see explained in the do file.


```
. lasso2 quality ${winevars} ///
>                      , plotpath(lambda) plotlabel ///
>                      plotvar(${winevars}) ///
>                      plotopt(legend(on)) alpha(1)
```

Knot	ID	Lambda	s	L1-Norm	EBIC	R-sq	Action
-----+-----							

1	1	1008.34183	1	0.00000	-677.11968	0.0000	Added _cons.
2	2	918.76352	2	0.15648	-711.69883	0.0259	Added volatileacidit
3	28	81.78971	3	1.61823	-925.16078	0.1516	Added fixedacidity.
4	39	29.39370	4	1.69937	-920.45671	0.1530	Added residualsugar.
5	59	4.57271	5	1.74881	-913.67292	0.1533	Added citricacid.

Use long option for full output.

Type e.g. `lasso2, lic(ebic)` to run the model selected by EBIC.



If we take a look at the full table, we notice that as lambda decreases, more variables are added to our model. (This is because the “penalty” or adding additional variables decreases as lambda decreases.) We can see this trend by looking at the generated output graph. As lambda increases, the coefficients for our covariates decrease until they reach 0 – that is, when they are dropped from the model all together. But notice how some variables, such as **citricacid** “drop off” quicker than others, such as **volatileacidity**. One might interpret this as an indication that variables that are “slower” to drop off are better predictors of wine quality; this may be a general rule of thumb but is not always true. The iteration highlighted with a “*” will indicate the model with the lowest error (as measured by ebic). 7 variables are included in this model. Let’s note the value of lambda that gave us our lowest “error”: 51.99287. This is our “selected lambda”

Our next command will run standard OLS with the 7 variable selected by the lasso algorithm. It won’t give us the full output we’re used to with **reg**, but it will at least give us the value of the coefficients:

```
. lasso2, lic(ebic)
```

Use `lambda=89.76411116338411` (selected by EBIC).

Selected	Lasso	Post-est OLS
volatileacidity	-1.6046319	-1.7614378
Partialled-out*		
_cons	6.4829802	6.5657455

Cross-Validation

Cross-validation is a means of testing out-of-sample fit. We want to develop a model that is valid not only for our data in the data, but, in our case, for all wine. We'll do this by dividing the dataset into k folds, or a certain number of equal divisions. The standard number is 10, so since our dataset has about 1,600 observations, each *fold* will have about 160 wine samples. The algorithm then takes a large majority of the folds and uses this subset to develop its lasso model as we did above. Then it uses the remaining folds to "test" the model. The assumption is that, with large enough datasets, a randomly-divided dataset will generate enough variance between the "training" and "test" divisions that the "test" portion can approximate the differences of out-of-sample data.

A side-note on training and test data

This assumption I just explained about gets at what I see as a fundamental problem of applying machine-learning techniques to development problems: heterogeneity in data availability. Let's use an invented example. Let's say I want to predict likelihood of forest fires based on satellite images. I collect my data using the best publicly available images from NASA or NOAA that I can find, develop my theory of change, and maybe use lasso to fine-tune my model. This is great, but suppose NASA has much better quality images of North America than anywhere else. Even if I train my model on images from around the world, can I really say that my "test" data is good enough to create valid model for terrain that has different tree species, different roof coverings, and other differences that aren't captured in the lesser-quality images?

Back to Cross-validation

Lassopack also includes a cross-validation command called **cvlasso** that works very similarly to **lasso2**. We'll run this command, and also capture the selected variables in a global so we can run a better OLS command than last time. Keep in mind that if you type

```
. help cvlasso
```

you'll see that you can adjust the number of k-folds in the options.

```
. cvlasso          quality ${winevars} ///
>                  , plotcv seed(123) lopt alpha(1) postest
```

K-fold cross-validation with 10 folds. Elastic net with alpha=1.

Fold	1	2	3	4	5	6	7	8	9	10		Lambda	MSPE	st. dev.
1												1008.3418	.65131479	.0223197
2												918.76352	.63627149	.02174724

3	837.14311	.62217786	.02076667	
4	762.77362	.61047734	.01994615	
5	695.01091	.60076361	.01925827	
6	633.26806	.59269932	.01868041	
7	577.01027	.5860044	.01819395	
8	525.75027	.58044634	.01778355	
9	479.04406	.57583211	.01743655	
10	436.48711	.57200145	.01714248	
11	397.7108	.56882131	.01689271	^
12	362.37927	.56618122	.01668009	
13	330.18649	.56398949	.01649866	
14	300.85363	.56216998	.01634351	
15	274.12663	.56065948	.01621052	
16	249.77397	.55940553	.01609626	
17	227.58474	.55836456	.01599787	
18	207.36674	.55750039	.01591295	
19	188.94485	.55678302	.01583949	
20	172.15951	.5561875	.0157758	
21	156.86533	.55569314	.01572045	
22	142.92985	.55528277	.01567224	
23	130.23236	.5549518	.01562854	
24	118.66287	.55472942	.01557949	
25	108.12119	.55456506	.01553216	
26	98.516001	.55444559	.0154905	
27	89.764111	.55437042	.01545617	
28	81.789715	.55430874	.01542484	
29	74.523742	.55426189	.01539515	
30	67.903258	.55423364	.01536928	
31	61.87092	.55421047	.01535176	
32	56.374478	.55419819	.01534066	*
33	51.366325	.55420873	.01533232	
34	46.803082	.55423424	.01532959	
35	42.645225	.55426779	.01532958	
36	38.856741	.5543072	.01532826	
37	35.404815	.55434925	.0153272	
38	32.259549	.55439109	.01532627	
39	29.3937	.55443185	.01532555	
40	26.782444	.55447149	.01532519	
41	24.403166	.55450642	.01532603	
42	22.235256	.55454039	.01532689	
43	20.259937	.55457298	.01532778	
44	18.4601	.55462753	.01534374	
45	16.820156	.55469167	.01536667	
46	15.325899	.55475153	.01538795	
47	13.964389	.55479173	.01541062	
48	12.723831	.55482881	.01542978	
49	11.593481	.5548688	.01544855	
50	10.563548	.5549268	.01547055	
51	9.6251114	.55498171	.0154907	
52	8.770043	.55503821	.01550849	
53	7.9909366	.5550937	.01552424	
54	7.2810438	.55514735	.01553824	
55	6.6342159	.55519693	.01555107	
56	6.0448505	.55524269	.01556283	
57	5.5078426	.55528486	.01557359	
58	5.0185411	.55532368	.01558345	
59	4.5727077	.55535938	.01559247	
60	4.166481	.55539215	.01560071	
61	3.7963423	.55541975	.01560787	
62	3.4590857	.55544561	.01561451	
63	3.1517901	.55546934	.01562057	

64	2.8717937	.55549109	.01562611
65	2.6166715	.55551102	.01563117
66	2.3842136	.55552927	.01563579
67	2.1724066	.55554597	.01564001
68	1.979416	.55556126	.01564387
69	1.8035701	.55557523	.01564739
70	1.6433459	.55558801	.0156506
71	1.4973556	.55559969	.01565353
72	1.3643347	.55561037	.0156562
73	1.2431309	.55562011	.01565864
74	1.1326946	.55562902	.01566087
75	1.0320691	.55563715	.0156629
76	.94038297	.55564457	.01566475
77	.85684194	.55565134	.01566643
78	.78072246	.55565752	.01566797
79	.71136523	.55566267	.01566922
80	.6481695	.55566738	.01567037
81	.59058791	.55567116	.01567174
82	.53812171	.55567459	.01567301
83	.49031646	.55567772	.01567417
84	.4467581	.55568057	.01567522
85	.40706934	.5556794	.01567693
86	.37090642	.5556814	.01567788
87	.33795612	.55568324	.01567874
88	.30793303	.55568491	.01567953
89	.2805771	.55568643	.01568025
90	.25565141	.55568782	.0156809
91	.23294004	.55568909	.01568149
92	.21224629	.55569025	.01568204
93	.19339092	.5556913	.01568253
94	.1762106	.55569226	.01568298
95	.16055654	.55569314	.01568339
96	.14629314	.55569394	.01568377
97	.13329686	.55569466	.01568411
98	.12145513	.55569533	.01568442
99	.11066539	.55569593	.0156847
100	.10083418	.55569648	.01568496

* lopt = the lambda that minimizes MSPE.

Run model: cvlasso, lopt


^ lse = largest lambda for which MSPE is within one standard error of the minimal MSPE.

Run model: cvlasso, lse

Estimate lasso with lambda=56.374 (lopt).

Selected	Lasso	Post-est OLS
+-----		
fixedacidity	0.0038595	0.0119222
volatileacidity	-1.6533470	-1.7317454
+-----		
Partialled-out*		
+-----		
_cons	6.4765833	6.4508846
+-----		

. global lassovars = e(selected)



Notice how the graph now is different: we see the natural log of lambda (indicated by a red line) that generated the model with the lowest error. Notice that the cross-validation method gave us a different ideal lambda value. Since we stored the variables the `cvlasso` command selected, we can run OLS with these variables with

```
. reg          quality          ${lassovars}
```

Source	SS	df	MS	Number of obs	=	1,599
Model	159.610433	2	79.8052164	F(2, 1596)	=	144.32
Residual	882.55467	1,596	.552979117	Prob > F	=	0.0000
Total	1042.1651	1,598	.6521684	R-squared	=	0.1532
				Adj R-squared	=	0.1521
				Root MSE	=	.74363

quality	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
fixedacidity	.0119222	.0110529	1.08	0.281	-.0097576	.033602
volatileacidity	-1.731745	.1074738	-16.11	0.000	-1.94255	-1.520941
_cons	6.450885	.1212101	53.22	0.000	6.213137	6.688632

Now that we have significance values, we can look and see which chemical components are strong predictors of quality.

Ridge

Recall that ridge operates very similarly to lasso. The main difference is that, due to squared constraint term, no covariates will ever be *completely* dropped from the model; “unimportant” covariates will only see their coefficients reduced to near-zero.

Let's use a new dataset to explore ridge. This one from Prof. Hastie for predicting rates of prostate cancer. (Btw, his website is here for more info and datasets: <http://web.stanford.edu/~hastie/pub.htm>)

```
. import delimited using "https://web.stanford.edu/~hastie/ElemStatLearn/dataset
> s/prostate.data", clear
```

We'll follow the steps above, except we'll tell stata to set $\alpha=0$ in the options, which indicates a ridge regression. **Lpsa** is our outcome variable of interest, and we'll group the explanatory variables in a global like last time. Let's run a lasso before ridge so we can compare.

```
. global          rhsvars lcavol lweight ///
>                                     age lbph svi lcp gleason pgg45

. lasso2          lpsa ${rhsvars} , plotpath(lambda) plotlabel ///
>                                     plotvar(${rhsvars}) ///
>                                     plotopt(legend(on)) alpha(1) ///
>                                     lic(ebic) postresults long
```


Knot	ID	Lambda	s	L1-Norm	EBIC	R-sq	Action
1	1	163.62492	1	0.00000	31.41226	0.0000	Added _cons.
	2	149.08894	2	0.06390	26.66962	0.0916	Added lcavol.
2	3	135.84429	2	0.12213	18.19047	0.1676	
	4	123.77625	2	0.17518	10.54017	0.2307	
	5	112.78031	2	0.22352	3.69568	0.2832	
	6	102.76122	2	0.26757	-2.37828	0.3267	
	7	93.63220	2	0.30770	-7.72701	0.3628	
	8	85.31417	2	0.34427	-12.40327	0.3928	
	9	77.73509	3	0.40800	-12.63533	0.4221	Added svi.
3	10	70.82932	3	0.48389	-17.25324	0.4490	
	11	64.53704	4	0.60174	-18.31145	0.4801	Added lweight.
4	12	58.80375	4	0.71293	-23.37859	0.5066	
	13	53.57979	4	0.81423	-27.79633	0.5285	
	14	48.81991	4	0.90654	-31.62332	0.5468	
	15	44.48288	4	0.99065	-34.91944	0.5619	
	16	40.53114	4	1.06728	-37.74368	0.5745	
	17	36.93047	4	1.13711	-40.15255	0.5849	
	18	33.64967	4	1.20073	-42.19891	0.5936	
	19	30.66032	4	1.25870	-43.93126	0.6008	
	20	27.93654	4	1.31152	-45.39336*	0.6067	
	21	25.45474	5	1.35340	-42.20238	0.6123	Added pgg45.
5	22	23.19341	6	1.39138	-38.93672	0.6175	Added lbph.
	23	21.13297	6	1.42643	-40.17769	0.6224	
	24	19.25558	6	1.45836	-41.22016	0.6264	
	25	17.54496	6	1.48746	-42.09423	0.6298	
	26	15.98632	6	1.51397	-42.82594	0.6325	
	27	14.56614	6	1.53812	-43.43763	0.6349	
	28	13.27212	6	1.56013	-43.94842	0.6368	
	29	12.09306	7	1.58269	-39.94418	0.6389	Added age.
6	30	11.01875	7	1.61022	-40.80902	0.6421	
	31	10.03987	7	1.63531	-41.53293	0.6448	
	32	9.14796	7	1.65817	-42.13807	0.6470	
	33	8.33528	7	1.67900	-42.64335	0.6488	
	34	7.59480	7	1.69798	-43.06485	0.6503	
	35	6.92010	8	1.71689	-38.84649	0.6516	Added gleason.
7	36	6.30533	8	1.73692	-39.15171	0.6527	

9	37	5.74518	8	1.75517	-39.40584	0.6536	Added lcp.
	38	5.23480	8	1.77180	-39.61733	0.6544	
	39	4.76975	8	1.78695	-39.79327	0.6550	
	40	4.34602	8	1.80076	-39.93957	0.6555	
	41	3.95993	9	1.83346	-35.69248	0.6567	
	42	3.60814	9	1.86831	-36.01479	0.6578	
	43	3.28761	9	1.90006	-36.28320	0.6588	
	44	2.99554	9	1.92900	-36.50659	0.6596	
	45	2.72943	9	1.95536	-36.69246	0.6602	
	46	2.48695	9	1.97938	-36.84703	0.6607	
	47	2.26602	9	2.00126	-36.97555	0.6612	
	48	2.06471	9	2.02120	-37.08238	0.6616	
	49	1.88129	9	2.03937	-37.17116	0.6619	
	50	1.71416	9	2.05593	-37.24493	0.6621	
	51	1.56188	9	2.07101	-37.30622	0.6623	
	52	1.42313	9	2.08476	-37.35713	0.6625	
	53	1.29670	9	2.09728	-37.39942	0.6627	
	54	1.18150	9	2.10869	-37.43454	0.6628	
	55	1.07654	9	2.11909	-37.46371	0.6629	
	56	0.98091	9	2.12856	-37.48793	0.6630	
	57	0.89376	9	2.13720	-37.50804	0.6630	
	58	0.81437	9	2.14506	-37.52475	0.6631	
	59	0.74202	9	2.15223	-37.53862	0.6632	
	60	0.67610	9	2.15876	-37.55013	0.6632	
	61	0.61604	9	2.16471	-37.55969	0.6632	
	62	0.56131	9	2.17013	-37.56763	0.6633	
	63	0.51145	9	2.17507	-37.57422	0.6633	
	64	0.46601	9	2.17957	-37.57970	0.6633	
	65	0.42461	9	2.18367	-37.58424	0.6633	
	66	0.38689	9	2.18741	-37.58801	0.6633	
	67	0.35252	9	2.19081	-37.59115	0.6633	
	68	0.32120	9	2.19391	-37.59375	0.6633	
	69	0.29267	9	2.19674	-37.59590	0.6634	
	70	0.26667	9	2.19932	-37.59770	0.6634	
	71	0.24298	9	2.20166	-37.59919	0.6634	
	72	0.22139	9	2.20380	-37.60042	0.6634	
	73	0.20172	9	2.20575	-37.60145	0.6634	
	74	0.18380	9	2.20752	-37.60230	0.6634	
	75	0.16748	9	2.20914	-37.60301	0.6634	
	76	0.15260	9	2.21062	-37.60359	0.6634	
	77	0.13904	9	2.21196	-37.60408	0.6634	
	78	0.12669	9	2.21318	-37.60448	0.6634	
	79	0.11543	9	2.21430	-37.60482	0.6634	
	80	0.10518	9	2.21531	-37.60510	0.6634	
	81	0.09584	9	2.21624	-37.60533	0.6634	
	82	0.08732	9	2.21708	-37.60552	0.6634	
	83	0.07956	9	2.21785	-37.60568	0.6634	
	84	0.07250	9	2.21855	-37.60581	0.6634	
	85	0.06606	9	2.21919	-37.60592	0.6634	
	86	0.06019	9	2.21977	-37.60602	0.6634	
	87	0.05484	9	2.22030	-37.60609	0.6634	
	88	0.04997	9	2.22078	-37.60615	0.6634	
	89	0.04553	9	2.22122	-37.60621	0.6634	
	90	0.04148	9	2.22162	-37.60625	0.6634	
	91	0.03780	9	2.22199	-37.60629	0.6634	
	92	0.03444	9	2.22232	-37.60632	0.6634	
	93	0.03138	9	2.22262	-37.60634	0.6634	
	94	0.02859	9	2.22290	-37.60636	0.6634	
	95	0.02605	9	2.22315	-37.60638	0.6634	
	96	0.02374	9	2.22338	-37.60639	0.6634	
	97	0.02163	9	2.22359	-37.60640	0.6634	

98	0.01971	9	2.22378	-37.60641	0.6634
99	0.01796	9	2.22395	-37.60642	0.6634
100	0.01636	9	2.22411	-37.60643	0.6634

Use $\lambda=27.93654463358689$ (selected by EBIC).

Plotting only supported for list of lambda values.
Plotting options ignored.



value, and that we have three covariates with that value. Now let's run the same selection of variables with ridge.

Knot	ID	Lambda	s	L1-Norm	EBIC	R-sq	Action
1	1	1.636e+05	7	0.00550	31.10766	0.0036	Added lcavol lweight lbph svi lcp gleason _cons.
	2	1.491e+05	7	0.00603	31.07807	0.0039	
	3	1.358e+05	7	0.00661	31.04562	0.0043	
	4	1.238e+05	7	0.00726	31.01004	0.0047	
	5	1.128e+05	7	0.00796	30.97101	0.0052	
	6	1.028e+05	7	0.00873	30.92823	0.0057	
	7	9.363e+04	7	0.00958	30.88132	0.0062	
	8	8.531e+04	7	0.01051	30.82989	0.0068	
	9	7.774e+04	7	0.01152	30.77351	0.0075	
	10	7.083e+04	7	0.01264	30.71173	0.0082	
	11	6.454e+04	7	0.01386	30.64401	0.0090	
	12	5.880e+04	7	0.01520	30.56982	0.0099	
	13	5.358e+04	7	0.01666	30.48854	0.0108	
2	14	4.882e+04	8	0.01837	30.37788	0.0121	Added age.
	15	4.448e+04	8	0.02014	30.27832	0.0132	
	16	4.053e+04	8	0.02207	30.16931	0.0145	
	17	3.693e+04	8	0.02419	30.04999	0.0159	
	18	3.365e+04	8	0.02651	29.91940	0.0174	
3	19	3.066e+04	9	0.02916	29.56364	0.0212	Added pgg45.
	20	2.794e+04	9	0.03194	29.38726	0.0232	
	21	2.545e+04	9	0.03498	29.19446	0.0254	
	22	2.319e+04	9	0.03831	28.98378	0.0277	
	23	2.113e+04	9	0.04195	28.75367	0.0303	
	24	1.926e+04	9	0.04592	28.50245	0.0332	
	25	1.754e+04	9	0.05025	28.22832	0.0362	
	26	1.599e+04	9	0.05498	27.92935	0.0396	
	27	1.457e+04	9	0.06014	27.60349	0.0432	
	28	1.327e+04	9	0.06575	27.24856	0.0471	
	29	1.209e+04	9	0.07187	26.86224	0.0514	
	30	1.102e+04	9	0.07853	26.44207	0.0560	
	31	1.004e+04	9	0.08577	25.98547	0.0610	
	32	9147.95888	9	0.09364	25.48976	0.0664	
	33	8335.27943	9	0.10217	24.95213	0.0723	
	34	7594.79618	9	0.11143	24.36965	0.0785	
	35	6920.09542	9	0.12146	23.73936	0.0853	
	36	6305.33321	9	0.13232	23.05820	0.0925	
	37	5745.18479	9	0.14405	22.32310	0.1003	
	38	5234.79841	9	0.15671	21.53101	0.1086	
	39	4769.75334	9	0.17035	20.67891	0.1174	
	40	4346.02160	9	0.18503	19.76390	0.1268	
	41	3959.93302	9	0.20080	18.78321	0.1368	
	42	3608.14349	9	0.21771	17.73433	0.1473	
	43	3287.60596	9	0.23581	16.61501	0.1585	
	44	2995.54411	9	0.25515	15.42341	0.1702	
	45	2729.42823	9	0.27576	14.15813	0.1825	
	46	2486.95335	9	0.29768	12.81831	0.1954	
	47	2266.01927	9	0.32095	11.40377	0.2087	
	48	2064.71236	9	0.34557	9.91503	0.2226	
	49	1881.28900	9	0.37157	8.35342	0.2369	
	50	1714.16047	9	0.39895	6.72119	0.2516	
	51	1561.87918	9	0.42769	5.02152	0.2667	
	52	1423.12614	9	0.45777	3.25861	0.2821	
	53	1296.69954	9	0.48918	1.43767	0.2977	
	54	1181.50432	9	0.52185	-0.43509	0.3135	
	55	1076.54274	9	0.55575	-2.35244	0.3294	

	56	980.90565	9	0.59079	-4.30628	0.3453	
	57	893.76469	9	0.62690	-6.28767	0.3611	
	58	814.36510	9	0.66398	-8.28701	0.3768	
	59	742.01915	9	0.70195	-10.29418	0.3923	
	60	676.10021	9	0.74068	-12.29875	0.4075	
	61	616.03733	9	0.78006	-14.29014	0.4223	
	62	561.31027	9	0.81997	-16.25788	0.4368	
	63	511.44502	9	0.86029	-18.19177	0.4509	
	64	466.00965	9	0.90088	-20.08212	0.4644	
	65	424.61064	9	0.94161	-21.91988	0.4775	
	66	386.88940	9	0.98236	-23.69681	0.4900	
	67	352.51921	9	1.02300	-25.40561	0.5020	
	68	321.20238	9	1.06340	-27.03996	0.5134	
	69	292.66764	9	1.10345	-28.59459	0.5242	
4	70	266.66786	8	1.14310	-30.06520	0.5345	Removed age.
5	71	242.97782	9	1.18294	-31.44891	0.5442	Added age.
	72	221.39234	9	1.22213	-32.74329	0.5533	
	73	201.72445	9	1.26063	-33.94720	0.5619	
	74	183.80381	9	1.29835	-35.06032	0.5700	
	75	167.47519	9	1.33521	-36.08310	0.5776	
	76	152.59715	9	1.37114	-37.01671	0.5847	
	77	139.04084	9	1.40608	-37.86296	0.5914	
	78	126.68884	9	1.43999	-38.62419	0.5976	
	79	115.43415	9	1.47280	-39.30319	0.6033	
	80	105.17930	9	1.50449	-39.90316	0.6087	
	81	95.83546	9	1.53503	-40.42759	0.6137	
	82	87.32170	9	1.56438	-40.88022	0.6183	
	83	79.56428	9	1.59254	-41.26496	0.6226	
	84	72.49601	9	1.61950	-41.58586	0.6266	
	85	66.05566	9	1.64525	-41.84705	0.6302	
	86	60.18746	9	1.66979	-42.05269	0.6336	
	87	54.84057	9	1.69314	-42.20694	0.6366	
	88	49.96869	9	1.71530	-42.31395	0.6395	
	89	45.52961	9	1.73628	-42.37780	0.6421	
	90	41.48488	9	1.75613	-42.40252*	0.6444	
	91	37.79948	9	1.77484	-42.39203	0.6465	
	92	34.44148	9	1.79247	-42.35016	0.6485	
	93	31.38179	9	1.80902	-42.28063	0.6502	
	94	28.59392	9	1.82770	-42.18699	0.6518	
	95	26.05372	9	1.85317	-42.07267	0.6532	
	96	23.73917	9	1.87753	-41.94094	0.6545	
	97	21.63025	9	1.90079	-41.79487	0.6556	
	98	19.70868	9	1.92293	-41.63735	0.6566	
	99	17.95782	9	1.94395	-41.47105	0.6575	
	100	16.36249	9	1.96387	-41.29846	0.6583	

*indicates minimum EBIC.

Use lambda=41.48488213793183 (selected by EBIC).

Selected	Ridge	Post-est OLS
lcavol	0.4092614	0.5643413
lweight	0.5623717	0.6220198
age	-0.0113673	-0.0212482
lbph	0.0731408	0.0967125
svi	0.6032902	0.7616733
lcp	0.0204348	-0.1060509
gleason	0.0734774	0.0492279
pgg45	0.0027821	0.0044575

```

Partialled-out*|
-----+-----
              _cons |      -0.0872316      0.1815609
-----+-----

```


Plotting only supported for list of lambda values.
Plotting options ignored.

```

. global                                lpsaridge = e(selected)

. graph export                          "output/lpsa-ridge-graph.png", replace
(file output/lpsa-ridge-graph.png written in PNG format)

```



Notice how the ridge graph coefficient paths all approach zero as lambda approaches infinity, but never actually reach it. Likewise, if we compare the two following regressions with the lasso- and ridge-selected variables, we notice that the ridge regression includes all variables.

```

. eststo lasso: reg lpsa ${lpsalasso}

```

Source	SS	df	MS	Number of obs	=	97
Model	81.3492223	3	27.1164074	F(3, 93)	=	54.15
Residual	46.5684363	93	.500735874	Prob > F	=	0.0000
Total	127.917659	96	1.33247561	R-squared	=	0.6359
				Adj R-squared	=	0.6242
				Root MSE	=	.70763

lpsa	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
lcavol	.5258519	.0748632	7.02	0.000	.3771884	.6745154
lweight	.6617699	.1756352	3.77	0.000	.3129933	1.010547
svi	.6656665	.2070898	3.21	0.002	.2544271	1.076906
_cons	-.7771568	.6229995	-1.25	0.215	-2.01431	.4599967

```
. eststo ridge: reg lpsa ${lpsaridge}
```

Source	SS	df	MS	Number of obs	=	97
				F(8, 88)	=	21.68
Model	84.8592398	8	10.607405	Prob > F	=	0.0000
Residual	43.0584188	88	.489300213	R-squared	=	0.6634
				Adj R-squared	=	0.6328
Total	127.917659	96	1.33247561	Root MSE	=	.6995

lpsa	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
lcavol	.5643413	.0878335	6.43	0.000	.3897908	.7388918
lweight	.6220198	.2008967	3.10	0.003	.2227799	1.02126
age	-.0212482	.0110841	-1.92	0.058	-.0432755	.0007791
lbph	.0967125	.0579127	1.67	0.098	-.0183768	.2118018
svi	.7616733	.2411757	3.16	0.002	.2823873	1.240959
lcp	-.1060509	.089868	-1.18	0.241	-.2846446	.0725427
gleason	.0492279	.1553407	0.32	0.752	-.259479	.3579349
pgg45	.0044575	.0043653	1.02	0.310	-.0042177	.0131327
_cons	.1815609	1.320568	0.14	0.891	-2.442791	2.805913